

Verity Insurance · 3-Minute Demo Runbook

Snowflake Summit · Pediment Booth · for the recorded video walkthrough of the Verity Insurance dbt-wizard demo. Looped unattended on the booth screen.

3-5 days

MANUAL MODEL BUILD TODAY

90 sec

DBT-WIZARD BUILD

\$8.4M

RESERVE GAP EXPOSURE

12 hr

EARLIER THAN RESERVES COMMITTEE

ODI architecture · five layers, one loop

Fivetran

Connectors + Custom SDK

MDLS

Managed Data Lake (Iceberg on S3)

dbt Labs

Bronze · Silver · Gold marts

dbt-wizard

Build-time AI · 90s authoring

Snowflake

Reads governed gold layer

DATA ENGINEERS ANCHOR HERE

Architecture · Wizard · Live · Pipeline

Show the loop, the four sub-agents, the materialization, the connector deep-links. This is your DE proof.

BUSINESS USERS ANCHOR HERE

Claims Radar, Policy Holdings, Cat-event Map, Underwriting

Show the answers, the KPI tiles, the segments, the executive read. The same governed gold layer powers all of it.

Frame the demo on business outcomes — never functionality

PAIN

The claims team can't size cat-event reserve gaps in time for the Reserves Committee

REASON

Every new severity cut = 3-5 day DE backlog

IMPACT

Under-reserved books, ratings downgrades, regulator scrutiny

CAPABILITY

Build-time AI on the governed gold layer

PROOF

Live model build in 90 seconds — on screen

The mantra: the win is not that dbt-wizard *can* author a model. The win is that the VP of Claims Analytics walks into the next Reserves Committee with a tested, lineage-tracked answer that did not exist this morning. Every beat below ties to that one sentence.

The 3-minute script · five beats

00:00 – 00:25 Beat 1 · The stakeholder hook

WHAT YOU SAY (25S)

"It is Friday morning at Verity Insurance. The VP of Claims Analytics needs to know: Why did Hurricane-Beryl claim severity for tier-2 commercial wind exposure exceed our model by 38%? The Reserves Committee happens in 12 hours. Today, an answer like this costs three-to-five days of data-engineering work and a missed meeting. We are about to change that."

WHAT YOU SHOW

Open the **Scenario page** (</scenario>). The VP of Claims Analytics's question is the first thing in frame. Let the T-minus counter tick.

OUTCOME → Anchors the audience on a real business question, not a feature.

00:25 – 00:50 Beat 2 · The cost of the gap

WHAT YOU SAY (25S)

"There is no gold-layer mart that answers this. Three days from now, after a request, a backlog, a scope, a build, and a peer review, the engineering team might land it. The Reserves Committee will have already happened. A \$8.4 million reserve gap that has to be sized before quarter-end — and that is just the dollars we can see."

WHAT YOU SHOW

Stay on Scenario. Highlight the four KPI tiles: 3-5 days, 90s, \$8.4M, T-minus to Reserves Committee.

OUTCOME → Quantifies the cost of the gap. Sets up the 90-second proof.

00:50 – 02:20 Beat 3 · The live build (proof moment · DE-anchor)

WHAT YOU SAY (90S)

"We click Run the Build. Four sub-agents — Explorer, Summary, Worker, Verification — open the dbt project, find the right grain, pull a slice of the gold layer, write the new model, write its tests, and materialize it. The viewer sees the SQL appear character by character. By the time we are done talking, the model exists, the tests pass, the lineage graph has updated."

While narrating, pause once at: "Worker is anchoring the CTE to the grain the VP of Claims Analytics asked for — not a notebook-grade approximation." Then again at: "Verification is writing a uniqueness test on the natural key. This is a governed gold asset, not a SQL snippet."

WHAT YOU SHOW

Click into [/wizard-live](#). Let the playback run at 1x. The agents type SQL and YAML in real time. The step rail moves Discovery → Schema → Inspection → Model Creation → Tests → Materialization. Terminal-window aesthetic.

OUTCOME → Data engineers see the loop end-to-end. The wow is that the output is a tested, governed asset — not a snippet pasted into Slack.

02:20 – 02:45 Beat 4 · The outcome (P&L tie-back · business-anchor)

WHAT YOU SAY (25S)

"The new mart is materialized. The Reserves Committee convenes in 12 hours, not three days. The VP of Claims Analytics walks in with the answer: tier-2 commercial wind policies in coastal Texas and Louisiana share an outdated debris-removal sublimit that has not been re-tuned since 2024 — 22 of 34 over-model claims share this clause signature. That is a sublimit re-tune, not a reinsurance restructure. And the next time anyone asks this question — flood, wildfire, hail, anywhere — the model is already there."

WHAT YOU SHOW

Cut to [/outcome](#). Materialized model, tests passed, lineage updated, root-cause panel, impact panel. Then quick scroll through the business-user pages (Claims Radar, Policy Holdings, Cat-event Map, Underwriting) so the audience sees the same gold layer powering the executive read.

OUTCOME → Closes the loop from question to action AND shows the business-user surface running on the same data.

02:45 – 03:00 Beat 5 · The architectural ask

WHAT YOU SAY (15S)

"Fivetran lands every source. MDLS owns the open lake. dbt governs the layer. dbt-wizard authors the gaps. Snowflake reads it. That loop closes for cat-event reserve sizing, claims fraud, underwriting drift, and regulatory exam prep — all on the same governed lake. Want to see how this maps to your gold layer? Stop by the Pediment booth or scan the QR."

WHAT YOU SHOW

Final card with the five-pillar loop, QR code to [fivetran.com/contact](#). Tag dbt Labs co-brand.

OUTCOME → Calls out the architectural pattern. Invites a discovery conversation.

Business outcomes to weave through — never lead with features

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
Reserves Committee gets the answer same-day	3-5 day DE backlog	"In time for the meeting, not three days late"

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
\$8.4M exposure sized in 90 seconds	Open gap window	"The cost of the gap closed before lunch"
Reusable governed mart for every future ask	One-off notebook scripts	"The next person who asks gets it for free"
CDO does not add headcount to keep up	Hiring DE backfill	"Your team stops being the bottleneck"
Auditor sees lineage + tests, not a query history	Manual evidence packs	"Examiner-ready by default"

Anti-patterns · what NOT to say

- **Don't say:** "It runs dbt show against the warehouse and authors a new model file." **Say:** "Ninety seconds from question to materialized answer."
- **Don't say:** "Cortex Analyst." Cortex is intentionally absent from this demo. The story is build-time AI on the governed lake.
- **Don't say:** "dbt-wizard generates SQL." Say: "dbt-wizard authors a governed asset — schema contract, tests, lineage, ownership."
- **Don't say:** "Look at the typing animation." The animation is the texture, not the value.
- **Don't open with the architecture diagram.** Open with the VP of Claims Analytics's question. Architecture is the closer.
- **Don't skip the business pages.** After the wizard finishes, scroll through Claims Radar, Policy Holdings, Cat-event Map, Underwriting so the audience sees the persona surface that runs on the same gold layer.

Discovery prompts to anchor before the demo (if presenting live)

BEFORE YOU PRESS PLAY

- "What questions does your VP of Claims Analytics ask that your team cannot answer same-day?"
- "When a stakeholder asks for a new cut of data, what is the average time from ask to merged dbt model?"
- "How many one-off SQL scripts are sitting in someone's Snowflake history that never made it back into the governed layer?"
- "Who is your Economic Buyer for new data investment? Likely audience: CUO, CRO, Chief Actuary, VP Claims."

AFTER THE DEMO LANDS

- "Which of those agent steps — Discovery, Schema, Inspection, Model Creation, Tests, Materialization — is where your team spends the most time today?"
- "If we sized a pilot on three of your highest-value recurring VP of Claims Analytics questions, what would success look like in 90 days?"
- "Who needs to be in a 30-minute architecture review to see how this lands on your gold layer? Audience: CUO, CRO, Chief Actuary, VP Claims."

Personas at the booth — split the time

DATA ENGINEER · 60% OF BOOTH TIME

Scenario → Live → Outcome → Architecture → Pipeline

Lead with the question, show the four sub-agents, then walk the architecture (Fivetran + MDLS + dbt Labs + dbt-wizard + Snowflake) and the Pipeline page with real Fivetran connector deep-links.

BUSINESS USER · 40% OF BOOTH TIME

Claims Radar, Policy Holdings, Cat-event Map, Underwriting

After the wizard finishes, drop into the dashboard pages. Show that the same governed gold layer powers the CUO, CRO, Chief Actuary, Claims Operations read. The build-time AI is invisible to them — they see the answer.

Required follow-up materials

- One-pager · Verity Insurance · at fivetran-jasonchletsos.github.io/Fivetran-Demo-Repository/one-pager/insurance/
- Live demo URL · fivetran-jasonchletsos.github.io/Insurance-ODI-Demo/
- 3-minute walkthrough video (this recording)
- Architecture diagram on the Architecture tab of the demo
- Cross-industry credibility · Pediment Bank

Booth-floor reminder: the demo is built to run unattended on the booth screen. If a prospect walks up mid-playback, do not restart. Pick up the narration where the playback is — that is also the point. The build finishes whether someone is watching or not. That is the entire pitch.